



SAPIENZA  
UNIVERSITÀ DI ROMA

DEPARTMENT OF COMPUTER SCIENCE

**TicketCloud**  
CLOUD COMPUTING

**Professor:**

Emiliano Casalicchio

**Students:**

Andrea Di Marco

Jemuel Espiritu Fuentes

Michele Granatiero

{dimarco.1835169, fuentes.1803530, granatiero.1812623}@studenti.uniroma1.it

---

Academic Year 2022/2023

# Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Implementation</b>	<b>3</b>
2.1	Presentation Tier . . . . .	3
2.1.1	Authentication . . . . .	3
2.1.2	Events . . . . .	4
2.1.3	Tickets . . . . .	5
2.2	Logic Tier . . . . .	6
2.3	Data Tier . . . . .	6
<b>3</b>	<b>Testing</b>	<b>8</b>
3.1	Load Testing . . . . .	8
3.1.1	Low Load . . . . .	9
3.1.2	Average Load . . . . .	12
3.1.3	Heavy Load . . . . .	15
3.2	Stress Testing . . . . .	17
3.3	Endurance Testing . . . . .	18
3.4	Spike Testing . . . . .	19
<b>4</b>	<b>Future developments</b>	<b>22</b>

# 1 Introduction

For our Cloud Computing course project we designed and implemented a ticketing web application using technologies from different cloud providers such us *Amazon*, *OpenAI* and *Auth0*.

TicketCloud, with the resources provided by *Sapienza University*, leverages various services of *AWS* to build and host a robust application readily available to users. This is a key aspect for a service of this caliber since the expected traffic can be unpredictable, with uneven spikes and booking requests. Priority was also given to an accommodating user experience thanks to the use of Angular with Material design.

To make the best use of our resources we populated our website with content given by state of the art AI text and image generators such as *ChatGPT*, *Dalle2* and *NightCafe* respectively for the description of the fictional events and their cover images.

We then performed a number of different performance tests on the architecture such as load tests, endurance tests, stress tests and spike tests by using *JMeter*. We finally collected client side and server side data with the tools provided by *JMeter* and *AWS CloudWatch*.

## 2 Implementation

The implementation of our system is based on a multi-tier **serverless** architecture developed following the best practices advised by *AWS Well Architected* to the best of what the resources given to us by *Sapienza University* allowed. We circumvented the aforementioned limitations by making use of cloud services outside of the AWS ecosystem such as *Auth0*.

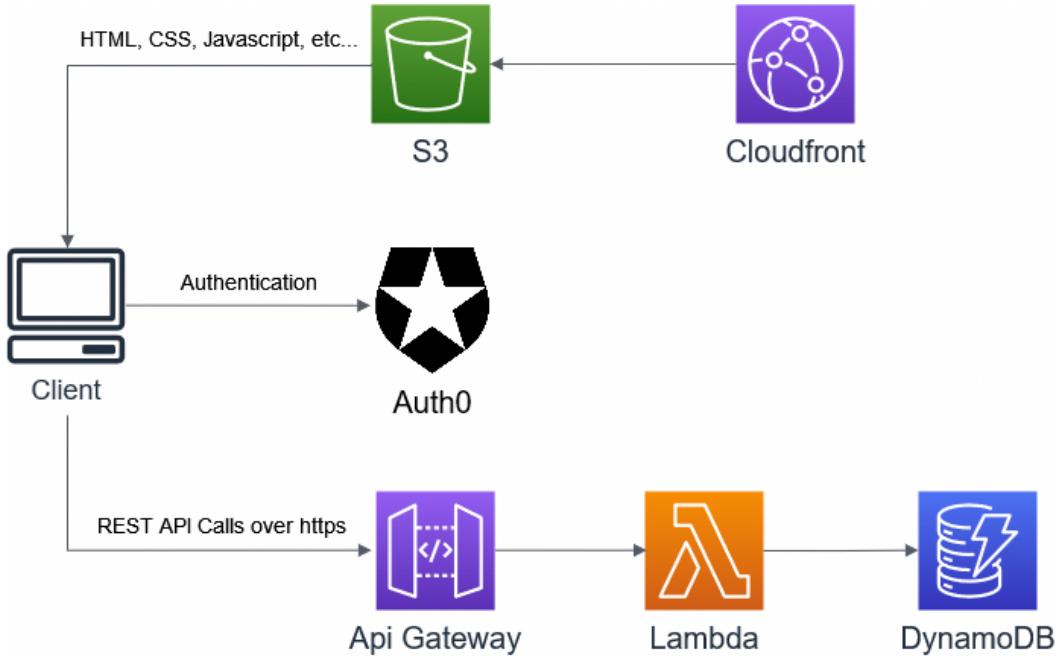


Figure 1: The architecture of TicketCloud

### 2.1 Presentation Tier

TicketCloud is built on the *Angular* framework and it is hosted on the AWS infrastructure using a combination of *AWS S3* and *AWS Cloudfront* as Content Delivery Network.

#### 2.1.1 Authentication

TicketCloud allows users to register and be recognized by implementing the authentication service provided by *Auth0*. The original plan was to use *AWS Cognito* as advised by *AWS Well Architected* but the *AWS Academy Learner Lab* account provided by *Sapienza University* wouldn't allow it.

*Auth0* takes care of the users' registration process by implementing connections with other platforms such as *Google*, *Facebook*, *GitHub* and more. Our users can also create brand new accounts that will then be required to verify by responding to an email sent to the provided address.

*Auth0* also allows fast login for users that have previously been authenticated and still have the cookies of our website in their browser.

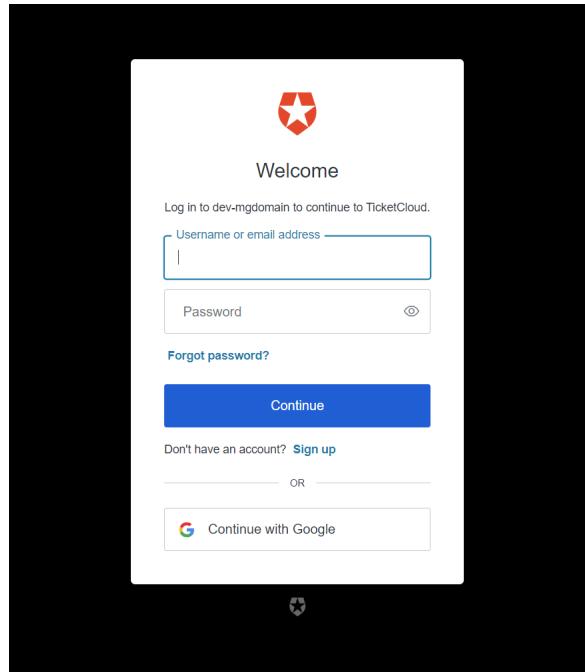


Figure 2: Authentication of website

### 2.1.2 Events

The landing page of TicketCloud serves the purpose of informing our users of all the available events. Said events are divided by the following categories:

- Music
- Show
- Other
- Sport

On top of the site we implemented an header containing the button to return to the dashboard and other useful functionalities such as the ability to search an event by beginning to type its name into a search bar and go to its respective page. The header also contains the authentication button thanks to which registration and log in, which make use of Auth0, can be carried out.

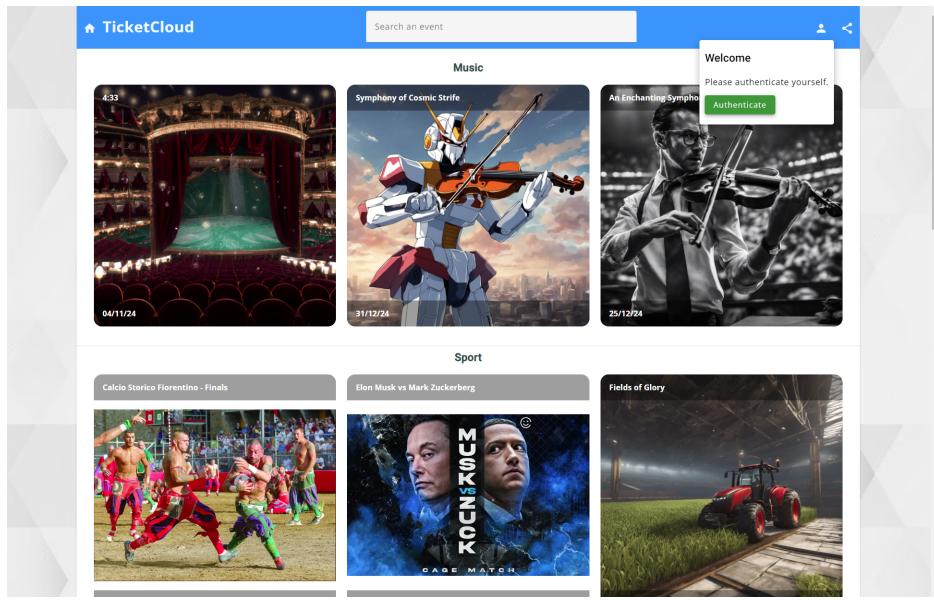


Figure 3: Dashboard of website

### 2.1.3 Tickets

Lastly we implemented the **event page** which contains every important information regarding the selected event. A user can read the description of the event, the date of the event and the number of available tickets.

Another important feature is the **purchase button**, which allows an **authenticated user** to book tickets for the selected event. This is composed of button and a menu that showcases the number of tickets the user wants. This button is disabled for guest users viewing the page.

After completing the procedure the user receives feedback from the system thanks to **Snackbar** component.

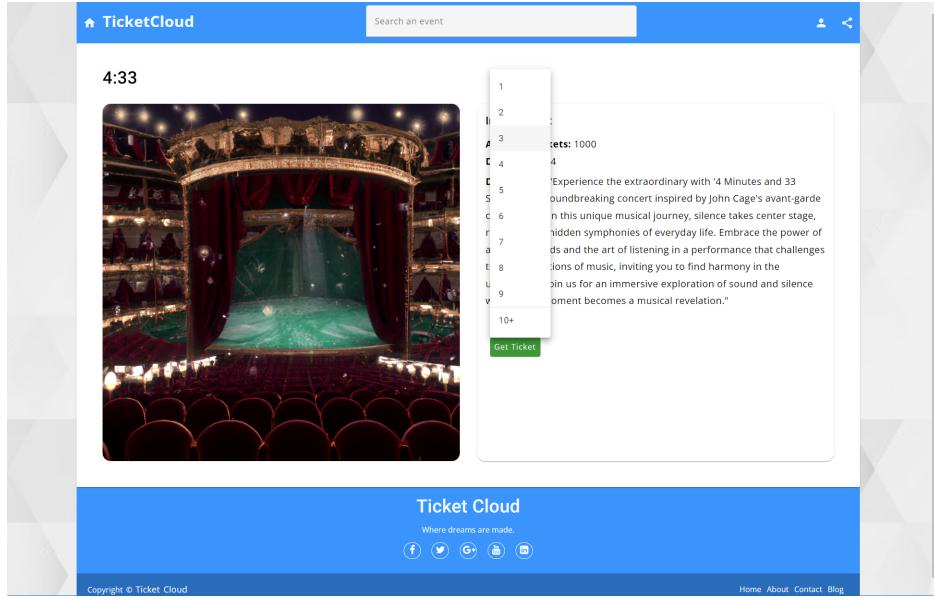


Figure 4: Event page and purchase button

## 2.2 Logic Tier

We decided to host the resources of the application on *AWS Simple Storage Service (S3)* and distribute it by using *AWS CloudFront*, the primary Content Delivery Network of *AWS* that also allows for SSL encryption.

To create and deploy our **RESTful APIs** we chose to use *AWS API Gateway* and linked the endpoints to our *Lambda* function thanks to Lambda Proxy Integration. We then added Cross-origin resource sharing (**CORS**) as an additional layer of security.

Our backend uses four *Lambda* functions to interact with the *DynamoDB* database:

- **getEvents**: queries the Event table and returns all the rows.
- **getEventById**: queries the Event table and returns the event with the given id.
- **getAvailableTickets**: queries the Event table and returns the number of available tickets for a given event
- **createTicket**: creates a specified number of entries in the Ticket table given an event **id**, a user email and the number of tickets to book while assigning a unique identifier to each entry

## 2.3 Data Tier

We chose to use *AWS DynamoDB* to store the data requested and created by our application. A ticketing service can be extremely heavy on its database usage as the number of users, events and tickets can be fairly high. *DynamoDB* allows us to have

a fast, flexible and reliable database for our usage that is also highly scalable. We decided to use 2 tables to store our data:

- Event

- **id** : number
- available\_tickets : number
- capacity : number
- description : string
- date : string
- name : string
- url : string

- Ticket

- **id** : string
- id\_event : number (Global Secondary Index)
- user\_email : string

A relevant aspect of the Ticket table is the way the single ticket's id is created. The *Lambda* function responsible for the ticket creation uses the timestamp and a *Universally Unique Identifier* generator to create a unique id across all tickets. This allows the system to have consistent data.

Another relevant aspect of the Ticket table is the use of a *Global Secondary Index* (GSI), which allows the table to define a second key to use for a query. This is an important optimization since we can search for every ticket for a given event without otherwise scanning the entire Ticket table, allowing us to **consume less resources**.

On the user side, *Auth0* takes care of the user data in its database. The informations are stored in the following table, accessible with the *Auth0 API*:

- User

- **email**
- username
- password
- connections

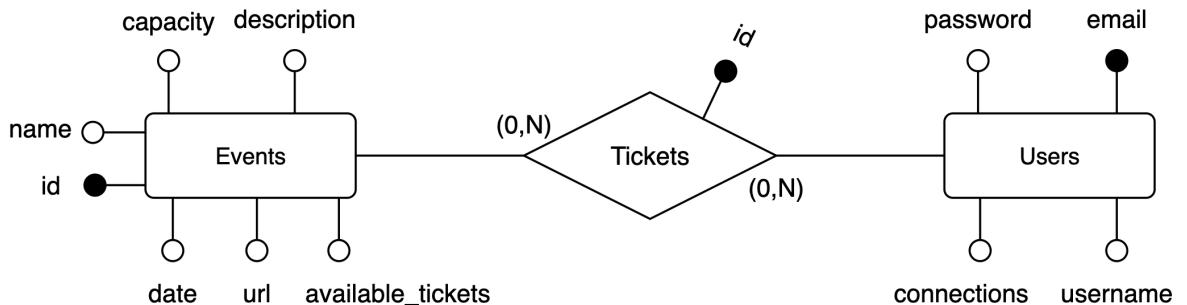


Figure 5: DB Entity-Relationship diagram.

## 3 Testing

In this section we show the results of our testing activities. The tool utilized is **JMeter**, an open source application for testing an analyzing the behaviour of a service. Every test is organized as follows:

1. **Warm up:** an initial stage that ensures consistency of the system
2. **Ramp up:** a stage in which the load is increased in a controlled manner
3. **Steady:** the main stage in which the load is constant representing the expected traffic
4. **Warm Down:** the final stage in which the load is gradually decreased to a minimum

Every phase is characterized by a duration and the amount of load which differ based on the test type. The duration of each phase was: 150 seconds, 600 seconds, 900 seconds and 150 seconds respectively and 30 minutes in total. The endurance test has double the time for each phase.



Figure 6: Expected user count through time.

### 3.1 Load Testing

We performed load testing tested the system with three different loads:

1. **Low Load** with 100 concurrent users.
2. **Average Load** with 500 concurrent users.
3. **Heavy Load** with 1000 concurrent users.



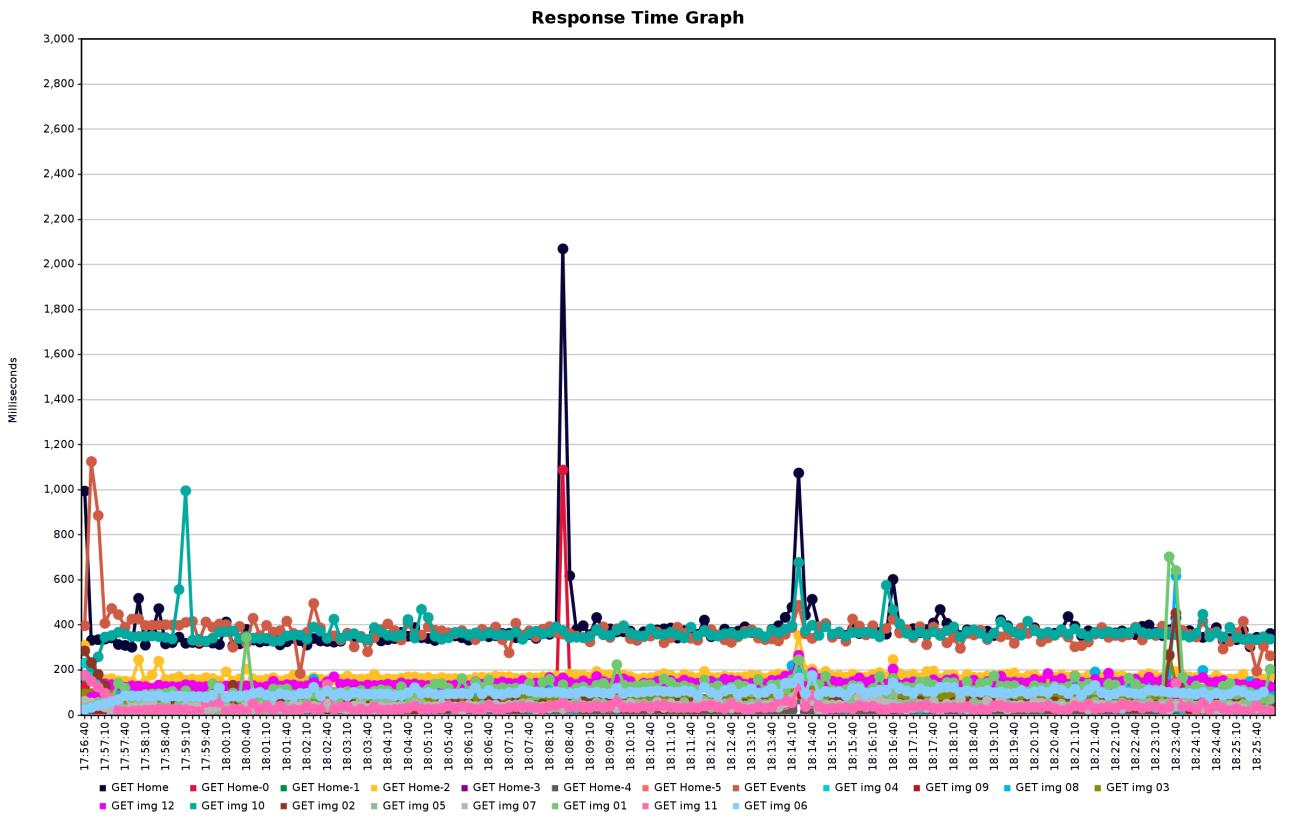


Figure 10: Response time graph of low load test landing

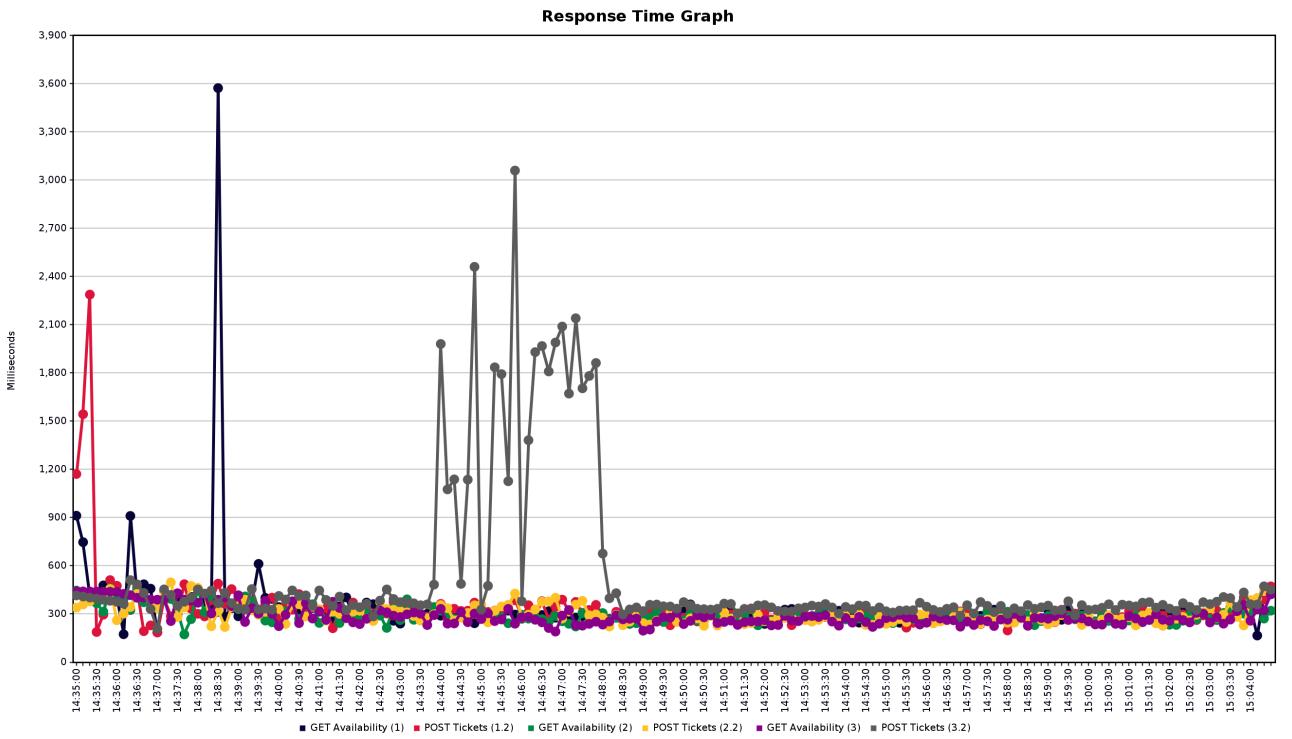


Figure 11: Response time graph of low load test booking

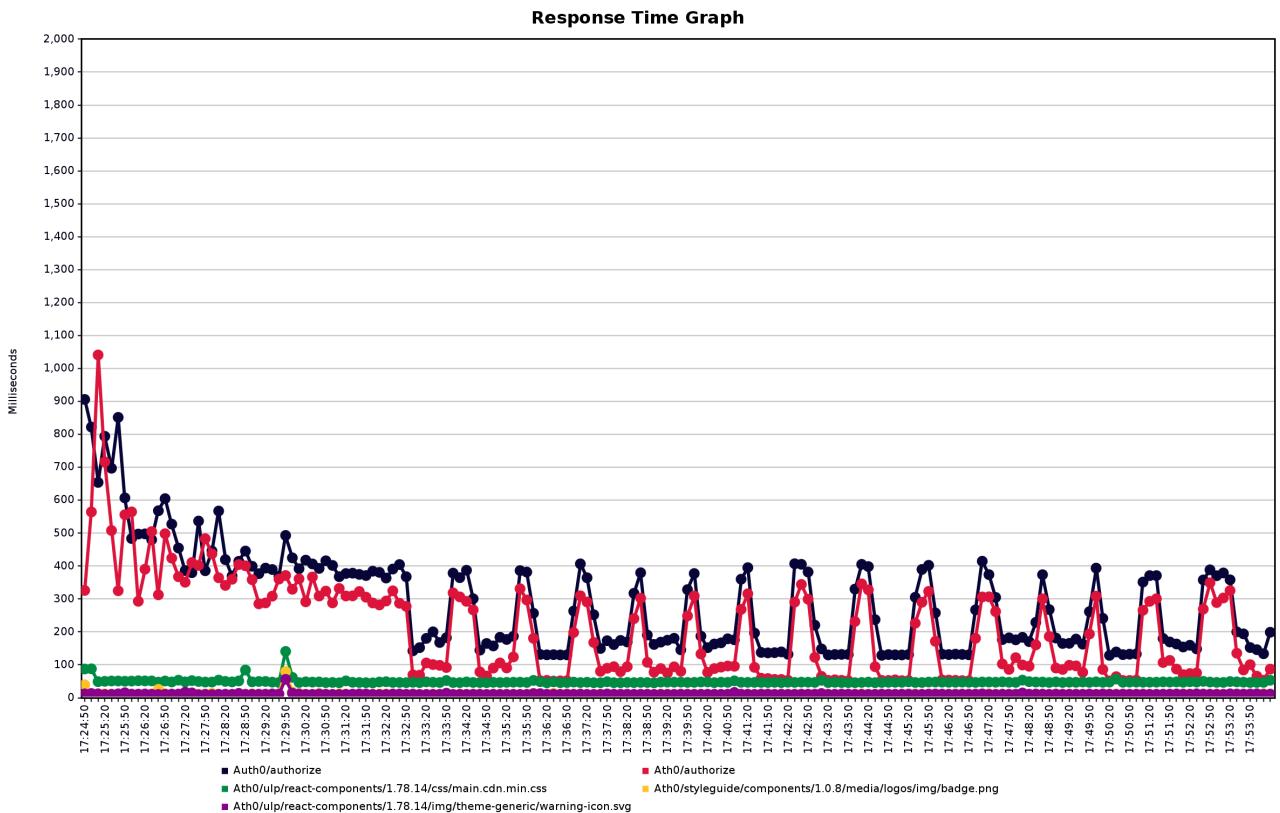


Figure 12: Response time graph of low load test login

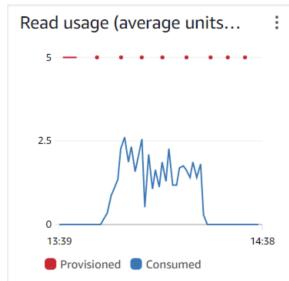


Figure 13: Read requests on Events

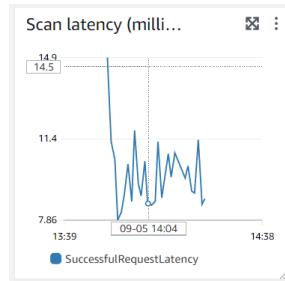


Figure 14: Latency on Events



Figure 15: Write requests on Tickets

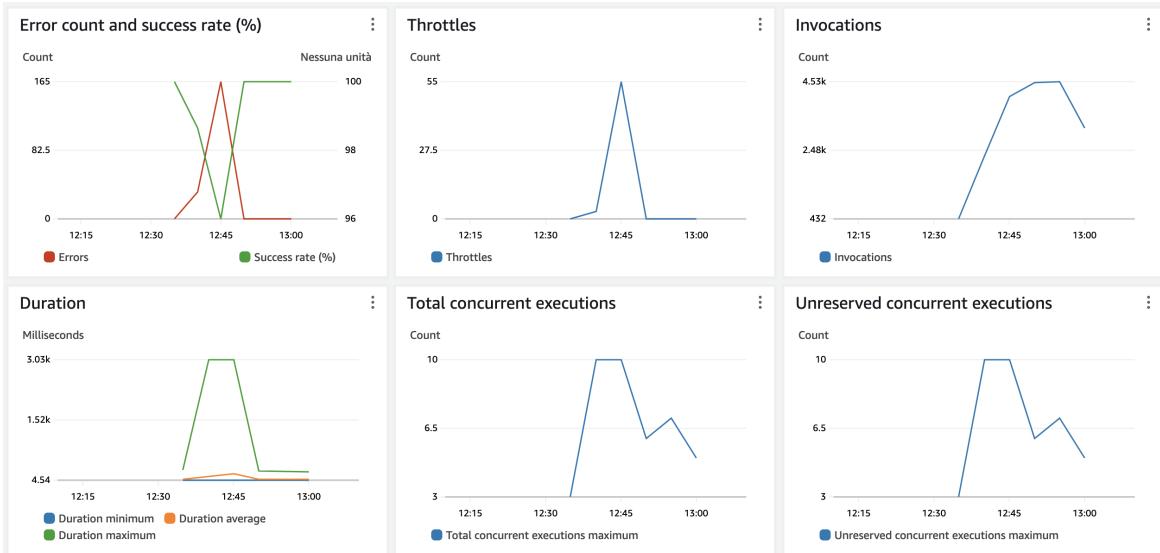


Figure 16: AWS lambda summary on low load

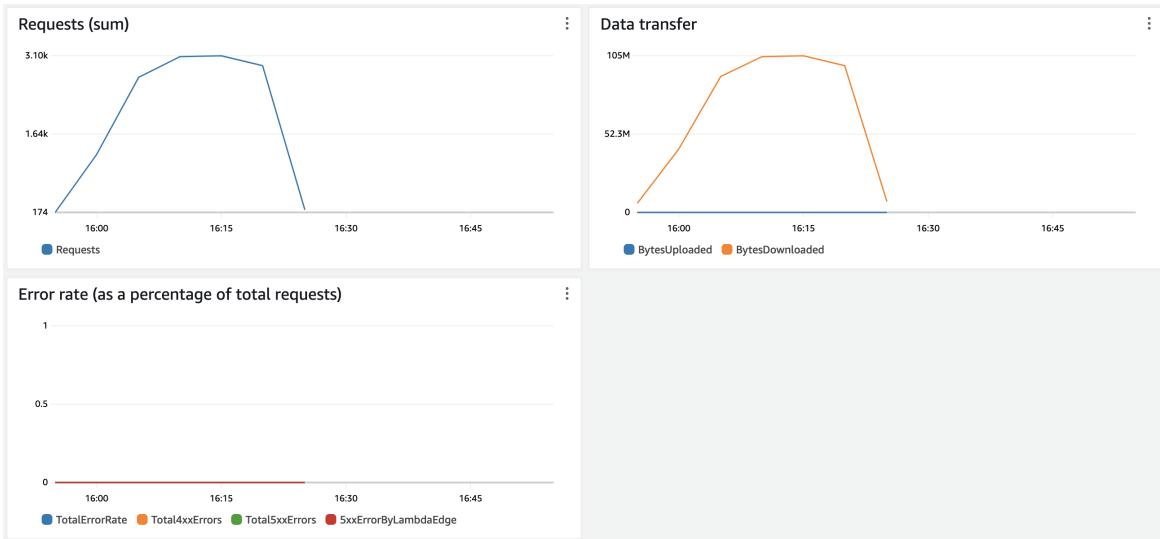


Figure 17: AWS Cloudfront on low load

### 3.1.2 Average Load

With a sustained load of 500 concurrent users we can see a significant increase in the response time during the **ramp up** phase for the booking procedure that it is then stabilized by the provisioning of more resources by *AWS* infrastructure. The landing procedure on the other hand struggled to handle the amount of concurrent users requesting the homepage. We can also see how the increase in load did not result in greater latency for *DynamoDB* operations although it did result in higher error rates for the Lambda functions, especially during the **ramp up** and **steady** phases. This was probably due to *AWS* not provisioning resources fast enough for the Academy Lab account.



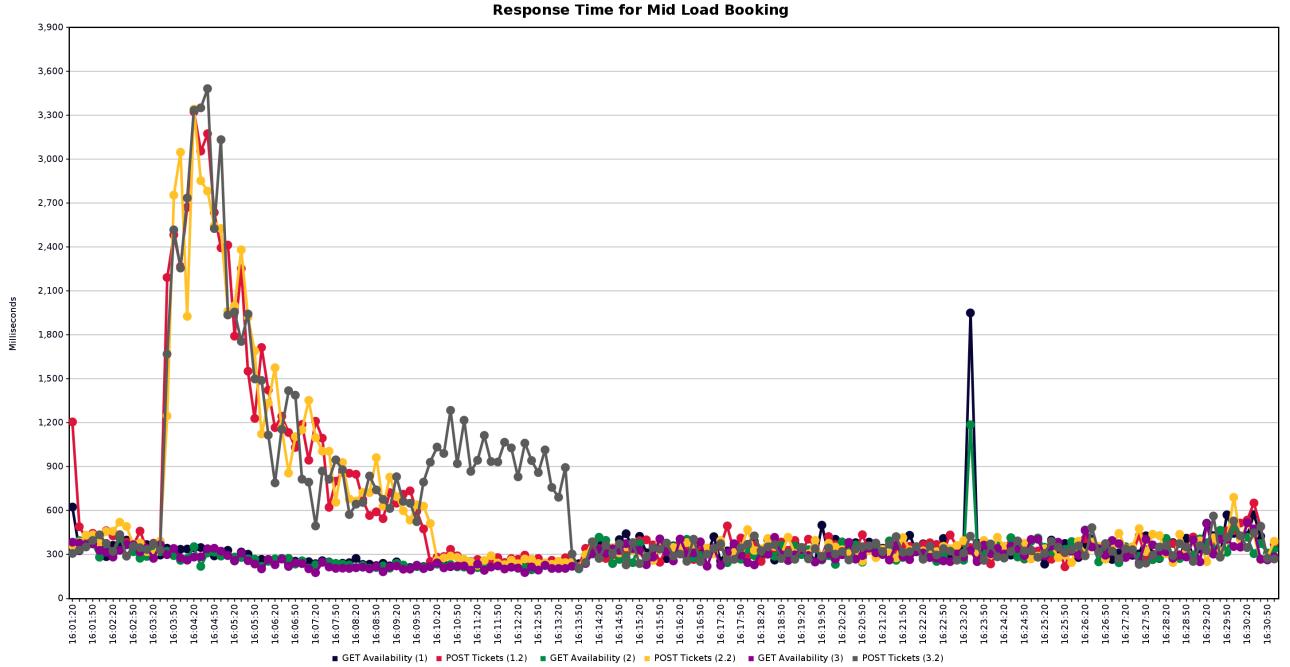


Figure 21: Response time graph of average load test booking

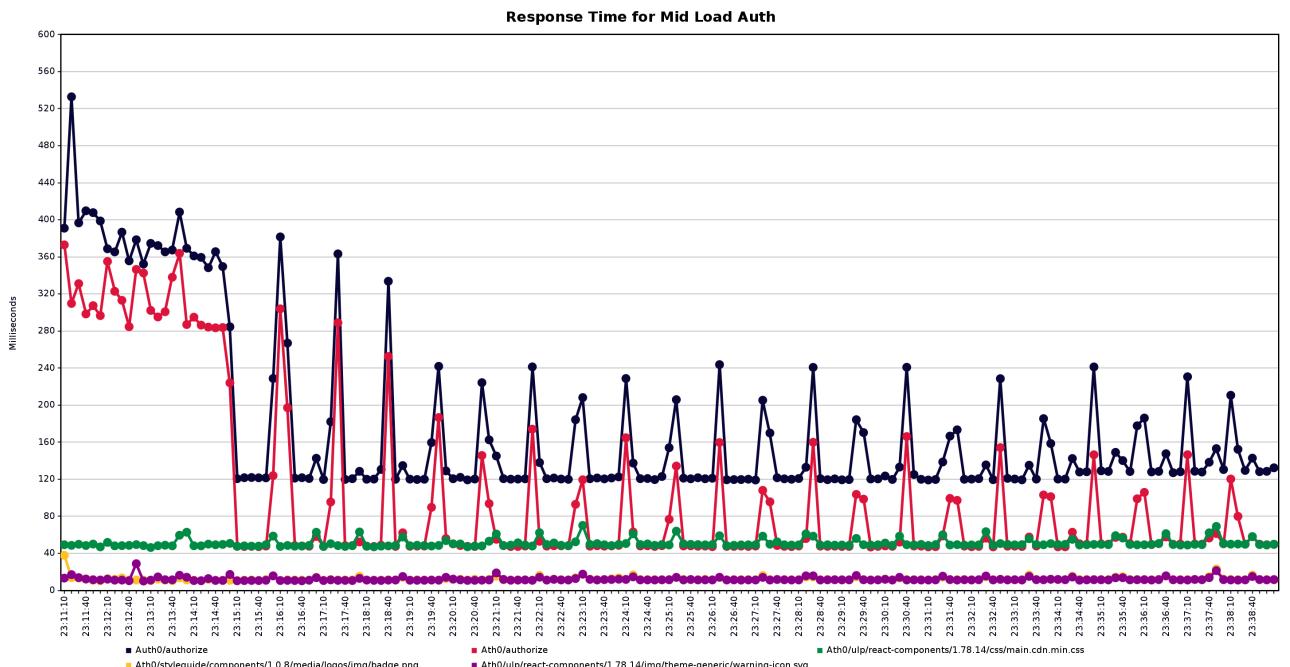


Figure 22: Response time graph of low load test login



Figure 23: Summary land AWS Dynamo events 14



Figure 24: Summary land AWS Dynamo events



Label ▲	# Samples	Average	Median	90% Line	95% Line	99% Line	Min	Maximum	Error %	Throughput	Received...	Sent KB/...
GET Availability (1)	32919	338	156	402	2361	4428	0	5509	87.96%	18.4/sec	31.33	1.76
GET Availability (2)	31800	325	153	386	1776	4177	0	5528	85.68%	17.9/sec	25.32	2.07
GET Availability (3)	31412	309	142	375	747	3396	0	5485	82.12%	17.8/sec	17.91	2.60
POST Tickets (1.2)	32759	394	156	512	3166	4690	0	10541	89.01%	18.3/sec	31.43	2.70
POST Tickets (2.2)	31613	382	154	442	3043	4480	0	10101	86.62%	17.8/sec	25.72	3.18
POST Tickets (3.2)	31220	392	144	433	3144	3600	0	10052	85.05%	17.8/sec	18.19	3.99
TOTAL	191723	357	154	396	2475	4260	0	10541	86.11%	106.9/sec	148.55	16.12

Figure 27: Summary of heavy load test booking

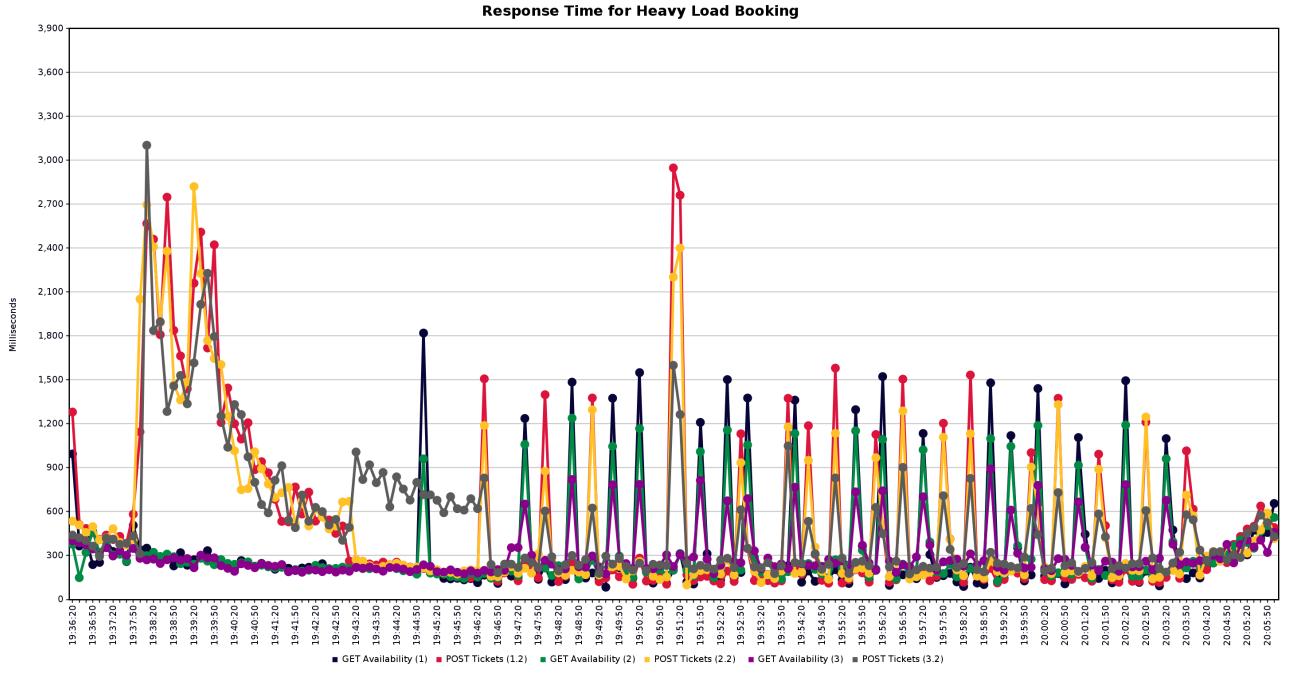


Figure 28: Response time graph of heavy load test booking

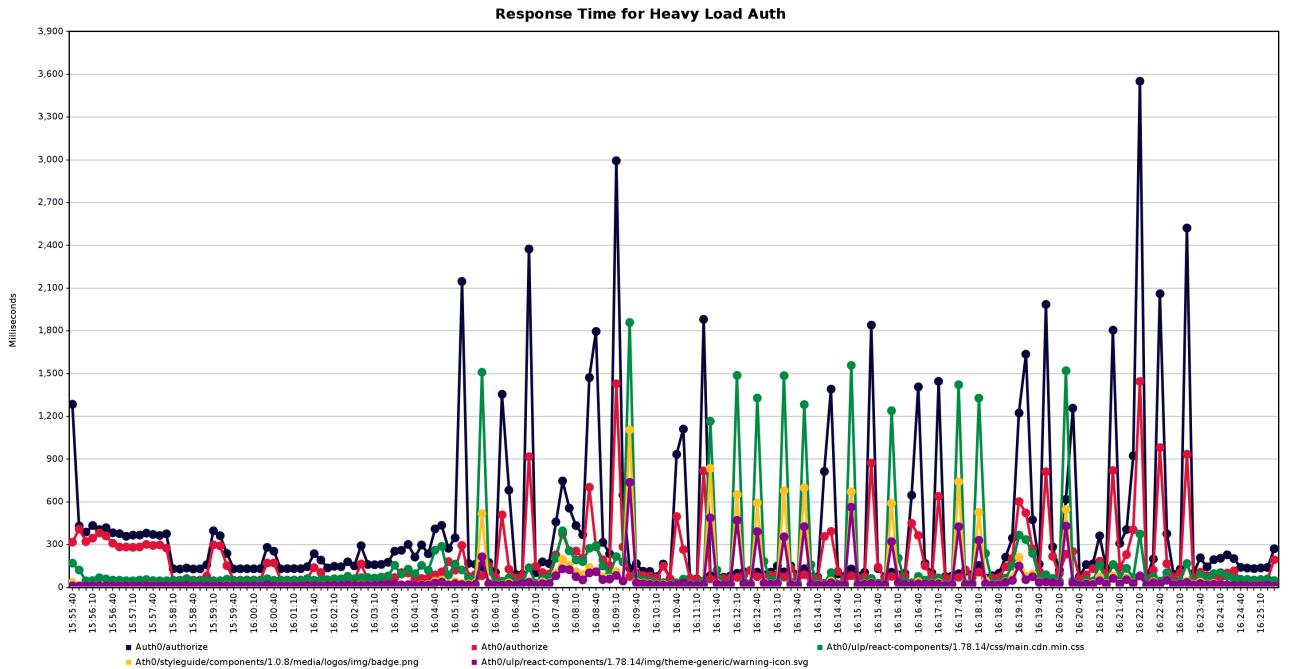


Figure 29: Response time graph of heavy load test login

### 3.2 Stress Testing

With a load of 5000 concurrent users response times are uniformly high and error rates are also higher as it is to be expected with a load this high for a small project. The purpose of this test were to learn the absolute limits our project could sustain.

*AWS CloudFront* managed to redirect landing calls efficiently and without fail even with such a stressful load. *AWS Lambda* returned a higher rate of failures again given to the concurrency limitations imposed by the *Academy Lab* account.

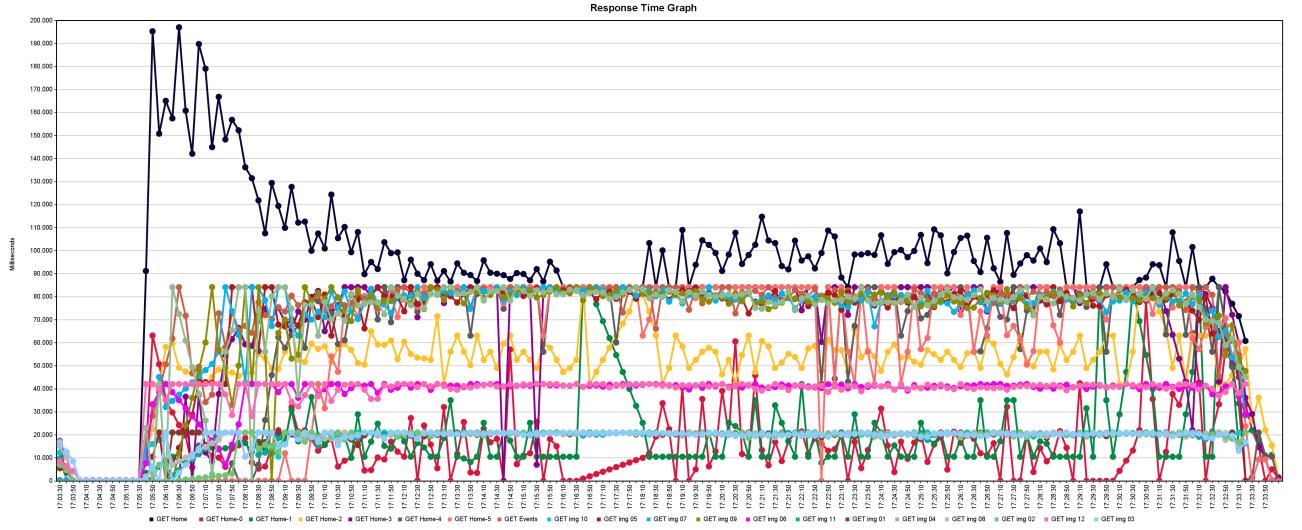


Figure 30: Response time graph of stress load test landing

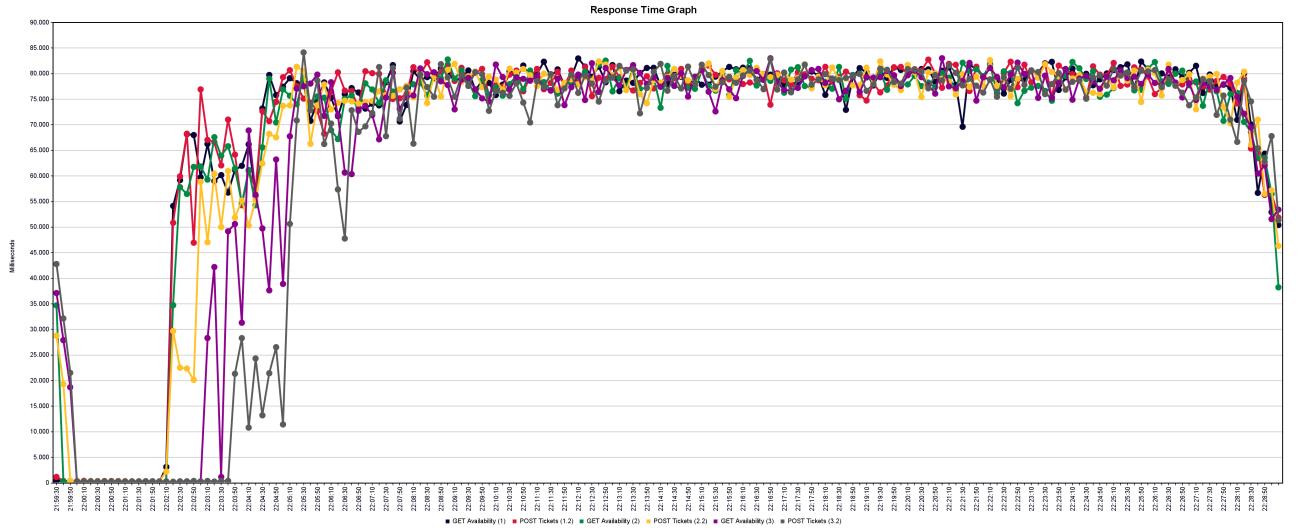


Figure 31: Response time graph of stress load test booking

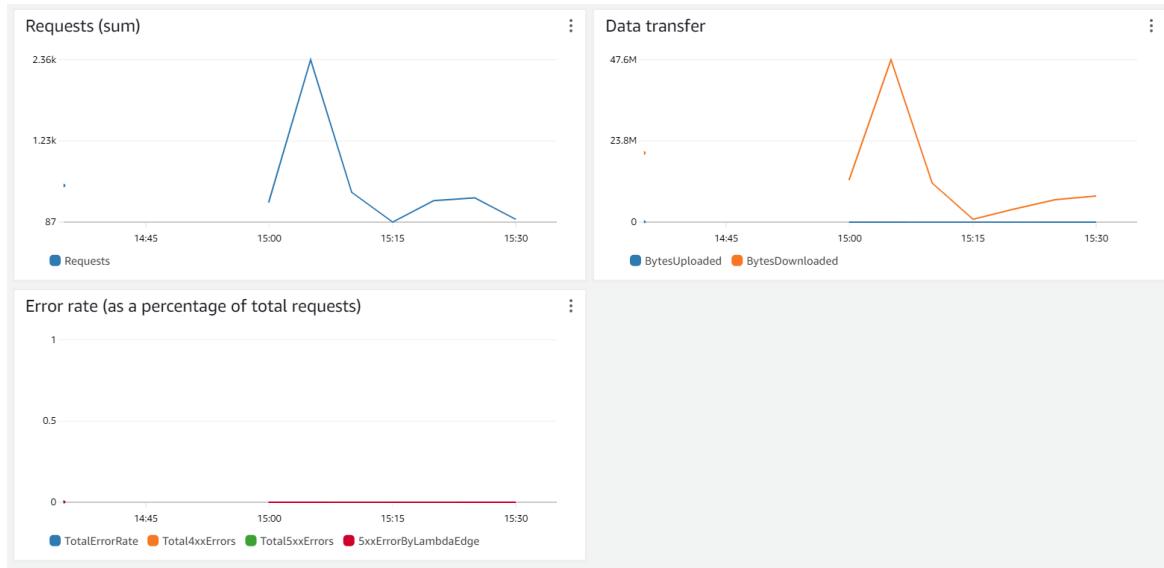


Figure 32: Summary stress AWS Cloudfront

### 3.3 Endurance Testing

For this test the system had to endure the expected load of 500 concurrent users for a full hour. Results show error rates and response time not too dissimilar from the average load test, with an increase during the ramp up phase and a subsequent decrease and stabilization for the other phases.

Label	# Samples	Average	Median	90% Line	95% Line	99% Line	Min	Maximum	Error %	Throughput	Received KB/s...	Sent KB/sec
GET Availability ...	32772	231	161	379	387	436	147	3412	74.77%	9.2/sec	4.72	1.65
POST Tickets ...	32663	492	162	699	3174	3466	146	4168	82.21%	9.1/sec	4.72	2.57
GET Availability ...	31706	224	157	374	381	416	127	3412	75.02%	8.9/sec	4.60	1.61
POST Tickets ...	31615	495	160	710	3170	3451	127	4483	82.66%	8.9/sec	4.59	2.50
GET Availability ...	31540	216	138	357	364	408	127	3417	75.20%	8.9/sec	4.59	1.61
POST Tickets ...	31441	521	161	871	3164	3438	127	3744	83.93%	8.9/sec	4.59	2.50
TOTAL	191737	363	158	388	795	3386	127	4483	78.95%	53.4/sec	27.58	12.34

Figure 33: Summary of endurance load test booking



Label ▲	# Samples	Average	Median	90% Line	95% Line	99% Line	Min	Maximum	Error %	Throughput	Received...	Sent KB...
GET Availability (1)	2476	413	160	430	2828	4767	1	5080	76.86%	17.2/sec	25.92	1.90
GET Availability (2)	2232	328	156	383	1370	4528	1	5290	80.65%	15.8/sec	20.19	1.98
GET Availability (3)	1882	273	139	363	379	3901	1	5225	87.94%	14.7/sec	13.87	2.20
POST Tickets (1.2)	2309	305	157	482	847	4257	1	5559	81.85%	16.2/sec	23.60	2.86
POST Tickets (2.2)	2062	523	156	2737	3376	4365	1	5862	92.39%	15.2/sec	18.84	3.05
POST Tickets (3.2)	1705	514	138	2553	3366	4083	1	5077	98.89%	13.3/sec	11.79	3.19
TOTAL	12666	389	156	448	3054	4467	1	5862	85.58%	85.3/sec	106.33	13.91

Figure 36: Summary of spike load test booking

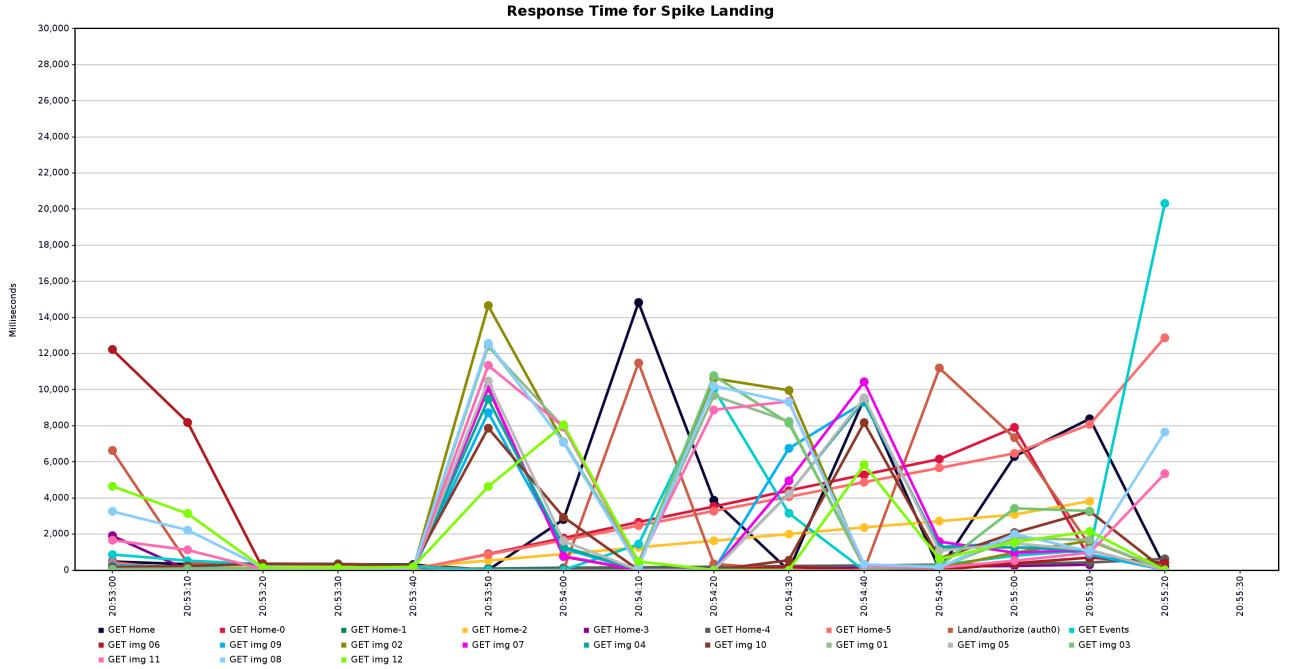


Figure 37: Response time graph of spike load test landing

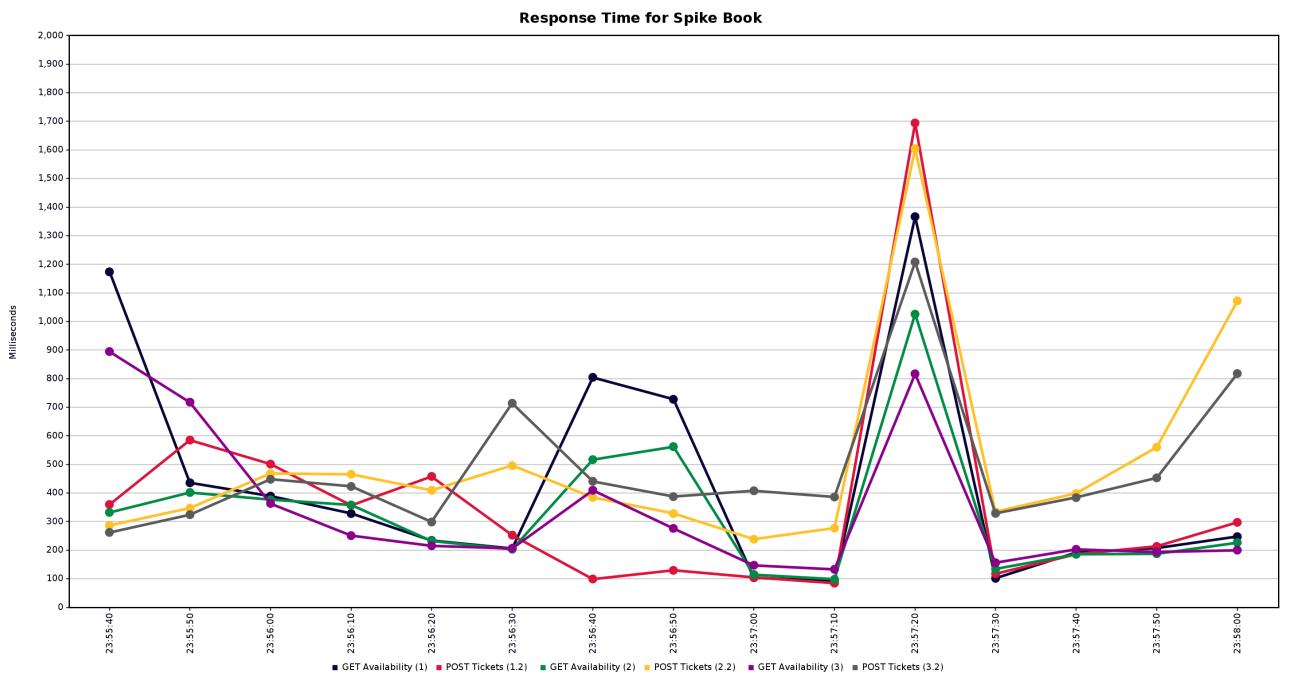


Figure 38: Response time graph of spike load test booking



Figure 39: Summary spike Dynamo tickets

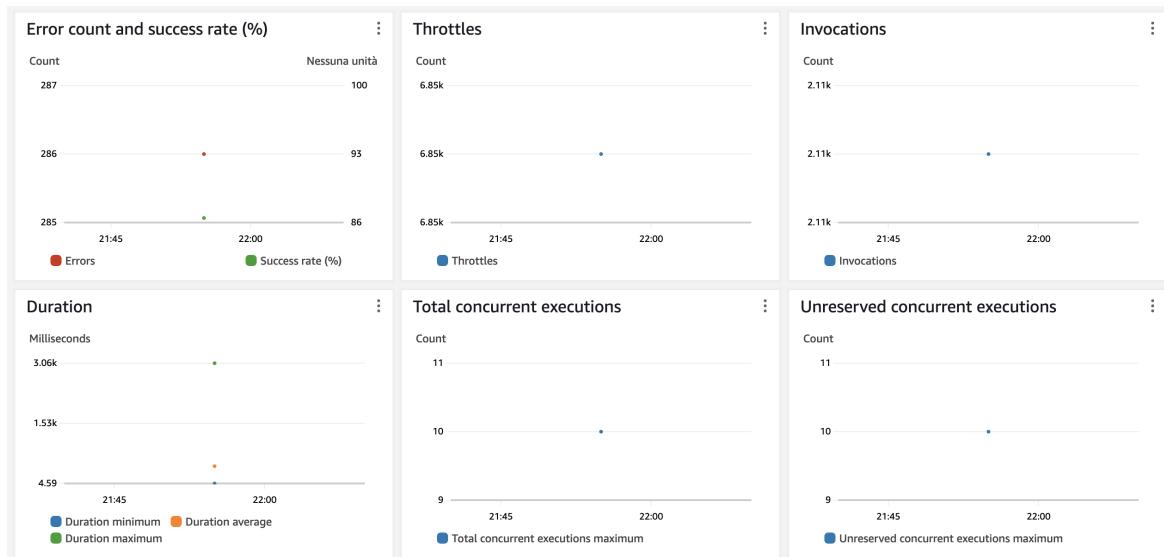


Figure 40: Summary spike AWS lambda

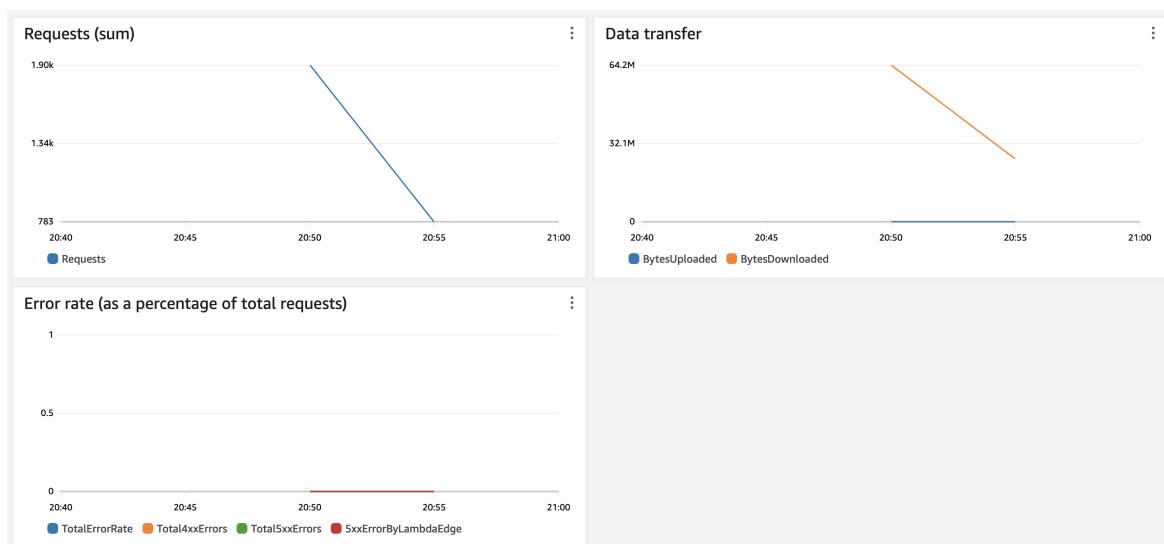


Figure 41: Summary spike AWS Cloudfront

## 4 Future developments

In this section we will propose some possible additions and activities to perform in the future in order to create a better platform for the user base.

First of all the limitations deriving from our *AWS Academy Lab* accounts meant that some useful services weren't available for use. Transitioning from a combo of *AWS S3* and *AWS Cloudfront* to a much simpler and more maintainable *AWS Amplify* solution would be an immediate option.

A second change in the architecture would be utilizing *AWS Cognito* instead of the external *Auth0* service for the authentication process. *Cognito* presents very useful features such as integration with other AWS services such as *Amplify*, *AWS IAM* and *AWS API Gateway*. In addition the cost of *Auth0* is greater.

Improving the event page would certainly be a priority and adding a functioning payment service would be the first step. A check against malicious users is also required to offer a truly professional ticketing service.

At the moment we have not developed a comprehensive user page but our data tier has the necessary information to showcase past purchases which could be useful. Additionally a recommendation system based on previous purchases is a normal service in today's applications and one should definitely be implemented in the future.

Regarding our architecture *DynamoDB* was our immediate choice for storing every information generated by our system and we believe it should be interesting to compare the performance relative to a similar system but with a traditional relational database such as *Amazon RDS*.

A similar activity could also be performed regarding the backend component. At the moment the infrastructure is completely serverless but it could be useful to analyze the behaviour of a serverful counterpart using services such as *AWS EC2* and *AWS ELB*.

Additional testing activities should also be considered such as testing for a longer period of time than 60 minutes and collecting a larger number of samples. Using AWS to create a distributed load testing platform could be a good option. Automating application testing would help with the future scaling of the system as a whole and it would be easily achieved by using services such as *EC2*, *Lambda* and *Amplify*.